

# Programming assignment 3: Optimization - Logistic Regression

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
```

## Your task

In this notebook code skeleton for performing logistic regression with gradient descent is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any `numpy` functions. No other libraries / imports are allowed.

For numerical reasons, we actually minimize the following loss function

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} NLL(\mathbf{w}) + \frac{1}{2} \lambda \|\mathbf{w}\|_2^2$$

where  $NLL(\mathbf{w})$  is the negative log-likelihood function, as defined in the lecture (see Slide 39).

## Load and preprocess the data

In this assignment we will work with the UCI ML Breast Cancer Wisconsin (Diagnostic) dataset <https://goo.gl/U2Uwz2>.

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image. There are 212 malignant examples and 357 benign examples.

```
In [3]: X, y = load_breast_cancer(return_X_y=True)

# Add a vector of ones to the data matrix to absorb the bias term
X = np.hstack([np.ones([X.shape[0], 1]), X])

# Set the random seed so that we have reproducible experiments
np.random.seed(123)

# Split into train and test
```

```
test_size = 0.3
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)
```

## Task 1: Implement the sigmoid function

```
In [4]: def sigmoid(t):
        """
        Applies the sigmoid function elementwise to the input data.

        Parameters
        -----
        t : array, arbitrary shape
            Input data.

        Returns
        -----
        t_sigmoid : array, arbitrary shape.
            Data after applying the sigmoid function.
        """

        # TODO
        return 1.0 / (1.0 + np.exp(-t))
```

## Task 2: Implement the negative log likelihood

As defined in Eq. 33

```
In [5]: def negative_log_likelihood(X, y, w):
        """
        Negative Log Likelihood of the Logistic Regression.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Classification targets.
        w : array, shape [D]
            Regression coefficients (w[0] is the bias term).

        Returns
        -----
        nll : float
            The negative log likelihood.
        """

        logit = sigmoid(np.dot(X, w))
        nll = -np.sum(y*np.log(logit) + (1-y)*np.log(1-logit))
        # TODO
        return nll
```

## Computing the loss function $\mathcal{L}(\mathbf{w})$ (nothing to do here)

```
In [6]: def compute_loss(X, y, w, lmbda):
        """
        Negative Log Likelihood of the Logistic Regression.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Classification targets.
        w : array, shape [D]
            Regression coefficients (w[0] is the bias term).
        lmbda : float
            L2 regularization strength.

        Returns
        -----
        loss : float
            Loss of the regularized logistic regression model.
        """
        # The bias term w[0] is not regularized by convention
        return negative_log_likelihood(X, y, w) / len(y) + lmbda * 0.5 * np.linalg.norm
```

## Task 3: Implement the gradient $\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w})$

Make sure that you compute the gradient of the loss function  $\mathcal{L}(\mathbf{w})$  (not simply the NLL!)

```
In [10]: def get_gradient(X, y, w, mini_batch_indices, lmbda):
        """
        Calculates the gradient (full or mini-batch) of the negative log likelihood w.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Classification targets.
        w : array, shape [D]
            Regression coefficients (w[0] is the bias term).
        mini_batch_indices: array, shape [mini_batch_size]
            The indices of the data points to be included in the (stochastic) calculation.
            This includes the full batch gradient as well, if mini_batch_indices = np.arange(N)
        lmbda: float
            Regularization strength. lmbda = 0 means having no regularization.

        Returns
        -----
        dw : array, shape [D]
            Gradient w.r.t. w.
        """
```

```

X_batch = X[mini_batch_indices]
y_batch = y[mini_batch_indices]

predictions = sigmoid(np.dot(X_batch, w))

dw = X_batch.T @ (predictions - y_batch) / len(mini_batch_indices)

dw[1:] += lambda * w[1:]
# TODO
return dw

```

## Train the logistic regression model (nothing to do here)

```

In [11]: def logistic_regression(X, y, num_steps, learning_rate, mini_batch_size, lambda, ver
        """
        Performs logistic regression with (stochastic) gradient descent.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Classification targets.
        num_steps : int
            Number of steps of gradient descent to perform.
        learning_rate: float
            The learning rate to use when updating the parameters w.
        mini_batch_size: int
            The number of examples in each mini-batch.
            If mini_batch_size=n_train we perform full batch gradient descent.
        lambda: float
            Regularization strength. lambda = 0 means having no regularization.
        verbose : bool
            Whether to print the loss during optimization.

        Returns
        -----
        w : array, shape [D]
            Optimal regression coefficients (w[0] is the bias term).
        trace: list
            Trace of the loss function after each step of gradient descent.
        """

        trace = [] # saves the value of loss every 50 iterations to be able to plot it
        n_train = X.shape[0] # number of training instances

        w = np.zeros(X.shape[1]) # initialize the parameters to zeros

        # run gradient descent for a given number of steps
        for step in range(num_steps):
            permuted_idx = np.random.permutation(n_train) # shuffle the data

            # go over each mini-batch and update the paramters
            # if mini_batch_size = n_train we perform full batch GD and this loop runs
            for idx in range(0, n_train, mini_batch_size):

```



Step 0, loss = 0.7427  
Step 50, loss = 0.9390  
Step 100, loss = 0.5168  
Step 150, loss = 0.3869  
Step 200, loss = 0.3676  
Step 250, loss = 0.3523  
Step 300, loss = 0.3396  
Step 350, loss = 0.3290  
Step 400, loss = 0.3198  
Step 450, loss = 0.3119  
Step 500, loss = 0.3050  
Step 550, loss = 0.2988  
Step 600, loss = 0.2934  
Step 650, loss = 0.2884  
Step 700, loss = 0.2840  
Step 750, loss = 0.2799  
Step 800, loss = 0.2763  
Step 850, loss = 0.2729  
Step 900, loss = 0.2698  
Step 950, loss = 0.2669  
Step 1000, loss = 0.2642  
Step 1050, loss = 0.2618  
Step 1100, loss = 0.2595  
Step 1150, loss = 0.2574  
Step 1200, loss = 0.2554  
Step 1250, loss = 0.2535  
Step 1300, loss = 0.2518  
Step 1350, loss = 0.2501  
Step 1400, loss = 0.2486  
Step 1450, loss = 0.2471  
Step 1500, loss = 0.2458  
Step 1550, loss = 0.2445  
Step 1600, loss = 0.2432  
Step 1650, loss = 0.2421  
Step 1700, loss = 0.2410  
Step 1750, loss = 0.2399  
Step 1800, loss = 0.2389  
Step 1850, loss = 0.2380  
Step 1900, loss = 0.2371  
Step 1950, loss = 0.2362  
Step 2000, loss = 0.2354  
Step 2050, loss = 0.2346  
Step 2100, loss = 0.2339  
Step 2150, loss = 0.2332  
Step 2200, loss = 0.2325  
Step 2250, loss = 0.2318  
Step 2300, loss = 0.2312  
Step 2350, loss = 0.2306  
Step 2400, loss = 0.2300  
Step 2450, loss = 0.2295  
Step 2500, loss = 0.2289  
Step 2550, loss = 0.2284  
Step 2600, loss = 0.2279  
Step 2650, loss = 0.2275  
Step 2700, loss = 0.2270  
Step 2750, loss = 0.2266

Step 2800, loss = 0.2261  
Step 2850, loss = 0.2257  
Step 2900, loss = 0.2253  
Step 2950, loss = 0.2249  
Step 3000, loss = 0.2246  
Step 3050, loss = 0.2242  
Step 3100, loss = 0.2239  
Step 3150, loss = 0.2235  
Step 3200, loss = 0.2232  
Step 3250, loss = 0.2229  
Step 3300, loss = 0.2226  
Step 3350, loss = 0.2223  
Step 3400, loss = 0.2220  
Step 3450, loss = 0.2217  
Step 3500, loss = 0.2214  
Step 3550, loss = 0.2212  
Step 3600, loss = 0.2209  
Step 3650, loss = 0.2206  
Step 3700, loss = 0.2204  
Step 3750, loss = 0.2202  
Step 3800, loss = 0.2199  
Step 3850, loss = 0.2197  
Step 3900, loss = 0.2195  
Step 3950, loss = 0.2193  
Step 4000, loss = 0.2191  
Step 4050, loss = 0.2189  
Step 4100, loss = 0.2187  
Step 4150, loss = 0.2185  
Step 4200, loss = 0.2183  
Step 4250, loss = 0.2181  
Step 4300, loss = 0.2179  
Step 4350, loss = 0.2177  
Step 4400, loss = 0.2175  
Step 4450, loss = 0.2174  
Step 4500, loss = 0.2172  
Step 4550, loss = 0.2170  
Step 4600, loss = 0.2169  
Step 4650, loss = 0.2167  
Step 4700, loss = 0.2166  
Step 4750, loss = 0.2164  
Step 4800, loss = 0.2163  
Step 4850, loss = 0.2161  
Step 4900, loss = 0.2160  
Step 4950, loss = 0.2158  
Step 5000, loss = 0.2157  
Step 5050, loss = 0.2156  
Step 5100, loss = 0.2154  
Step 5150, loss = 0.2153  
Step 5200, loss = 0.2152  
Step 5250, loss = 0.2151  
Step 5300, loss = 0.2149  
Step 5350, loss = 0.2148  
Step 5400, loss = 0.2147  
Step 5450, loss = 0.2146  
Step 5500, loss = 0.2145  
Step 5550, loss = 0.2144

Step 5600,	loss = 0.2142
Step 5650,	loss = 0.2141
Step 5700,	loss = 0.2140
Step 5750,	loss = 0.2139
Step 5800,	loss = 0.2138
Step 5850,	loss = 0.2137
Step 5900,	loss = 0.2136
Step 5950,	loss = 0.2135
Step 6000,	loss = 0.2134
Step 6050,	loss = 0.2133
Step 6100,	loss = 0.2132
Step 6150,	loss = 0.2131
Step 6200,	loss = 0.2130
Step 6250,	loss = 0.2129
Step 6300,	loss = 0.2128
Step 6350,	loss = 0.2128
Step 6400,	loss = 0.2127
Step 6450,	loss = 0.2126
Step 6500,	loss = 0.2125
Step 6550,	loss = 0.2124
Step 6600,	loss = 0.2123
Step 6650,	loss = 0.2122
Step 6700,	loss = 0.2122
Step 6750,	loss = 0.2121
Step 6800,	loss = 0.2120
Step 6850,	loss = 0.2119
Step 6900,	loss = 0.2118
Step 6950,	loss = 0.2118
Step 7000,	loss = 0.2117
Step 7050,	loss = 0.2116
Step 7100,	loss = 0.2115
Step 7150,	loss = 0.2114
Step 7200,	loss = 0.2114
Step 7250,	loss = 0.2113
Step 7300,	loss = 0.2112
Step 7350,	loss = 0.2112
Step 7400,	loss = 0.2111
Step 7450,	loss = 0.2110
Step 7500,	loss = 0.2109
Step 7550,	loss = 0.2109
Step 7600,	loss = 0.2108
Step 7650,	loss = 0.2107
Step 7700,	loss = 0.2107
Step 7750,	loss = 0.2106
Step 7800,	loss = 0.2105
Step 7850,	loss = 0.2105
Step 7900,	loss = 0.2104
Step 7950,	loss = 0.2103

[illegible]



```
lmbda=0.1,  
verbose=verbose)
```

Step 0, loss = 1.3392  
Step 50, loss = 0.3214  
Step 100, loss = 0.2859  
Step 150, loss = 0.2555  
Step 200, loss = 0.2583  
Step 250, loss = 0.2407  
Step 300, loss = 0.2287  
Step 350, loss = 0.2272  
Step 400, loss = 0.2222  
Step 450, loss = 0.2237  
Step 500, loss = 0.2221  
Step 550, loss = 0.2309  
Step 600, loss = 0.2157  
Step 650, loss = 0.2168  
Step 700, loss = 0.2145  
Step 750, loss = 0.2180  
Step 800, loss = 0.2122  
Step 850, loss = 0.2342  
Step 900, loss = 0.2111  
Step 950, loss = 0.2140  
Step 1000, loss = 0.2105  
Step 1050, loss = 0.2163  
Step 1100, loss = 0.2093  
Step 1150, loss = 0.2086  
Step 1200, loss = 0.2091  
Step 1250, loss = 0.2114  
Step 1300, loss = 0.2101  
Step 1350, loss = 0.2071  
Step 1400, loss = 0.2078  
Step 1450, loss = 0.2064  
Step 1500, loss = 0.2061  
Step 1550, loss = 0.2092  
Step 1600, loss = 0.2166  
Step 1650, loss = 0.2065  
Step 1700, loss = 0.2134  
Step 1750, loss = 0.2070  
Step 1800, loss = 0.2049  
Step 1850, loss = 0.2100  
Step 1900, loss = 0.2039  
Step 1950, loss = 0.2248  
Step 2000, loss = 0.2060  
Step 2050, loss = 0.2180  
Step 2100, loss = 0.2028  
Step 2150, loss = 0.2052  
Step 2200, loss = 0.2074  
Step 2250, loss = 0.2080  
Step 2300, loss = 0.2019  
Step 2350, loss = 0.2040  
Step 2400, loss = 0.2127  
Step 2450, loss = 0.2060  
Step 2500, loss = 0.2016  
Step 2550, loss = 0.2013  
Step 2600, loss = 0.2126  
Step 2650, loss = 0.2024  
Step 2700, loss = 0.2050  
Step 2750, loss = 0.2064

Step 2800, loss = 0.2000  
Step 2850, loss = 0.2074  
Step 2900, loss = 0.1997  
Step 2950, loss = 0.2157  
Step 3000, loss = 0.1997  
Step 3050, loss = 0.1998  
Step 3100, loss = 0.2021  
Step 3150, loss = 0.2119  
Step 3200, loss = 0.2046  
Step 3250, loss = 0.1985  
Step 3300, loss = 0.1983  
Step 3350, loss = 0.2023  
Step 3400, loss = 0.1994  
Step 3450, loss = 0.2036  
Step 3500, loss = 0.2076  
Step 3550, loss = 0.2109  
Step 3600, loss = 0.1975  
Step 3650, loss = 0.2200  
Step 3700, loss = 0.1985  
Step 3750, loss = 0.1984  
Step 3800, loss = 0.1980  
Step 3850, loss = 0.2021  
Step 3900, loss = 0.1988  
Step 3950, loss = 0.2152  
Step 4000, loss = 0.2016  
Step 4050, loss = 0.1963  
Step 4100, loss = 0.2036  
Step 4150, loss = 0.1965  
Step 4200, loss = 0.1978  
Step 4250, loss = 0.2101  
Step 4300, loss = 0.1964  
Step 4350, loss = 0.1957  
Step 4400, loss = 0.2229  
Step 4450, loss = 0.1960  
Step 4500, loss = 0.2022  
Step 4550, loss = 0.1953  
Step 4600, loss = 0.1955  
Step 4650, loss = 0.1952  
Step 4700, loss = 0.2042  
Step 4750, loss = 0.1973  
Step 4800, loss = 0.1973  
Step 4850, loss = 0.1949  
Step 4900, loss = 0.1954  
Step 4950, loss = 0.1953  
Step 5000, loss = 0.1981  
Step 5050, loss = 0.1944  
Step 5100, loss = 0.2009  
Step 5150, loss = 0.1944  
Step 5200, loss = 0.1987  
Step 5250, loss = 0.2079  
Step 5300, loss = 0.2014  
Step 5350, loss = 0.1951  
Step 5400, loss = 0.1960  
Step 5450, loss = 0.1981  
Step 5500, loss = 0.1938  
Step 5550, loss = 0.1990

Step 5600, loss = 0.1949  
Step 5650, loss = 0.1983  
Step 5700, loss = 0.1938  
Step 5750, loss = 0.1937  
Step 5800, loss = 0.1939  
Step 5850, loss = 0.2032  
Step 5900, loss = 0.1937  
Step 5950, loss = 0.1948  
Step 6000, loss = 0.1932  
Step 6050, loss = 0.1929  
Step 6100, loss = 0.1952  
Step 6150, loss = 0.1928  
Step 6200, loss = 0.2002  
Step 6250, loss = 0.1925  
Step 6300, loss = 0.1925  
Step 6350, loss = 0.1985  
Step 6400, loss = 0.1944  
Step 6450, loss = 0.1958  
Step 6500, loss = 0.1930  
Step 6550, loss = 0.1937  
Step 6600, loss = 0.1921  
Step 6650, loss = 0.2048  
Step 6700, loss = 0.1996  
Step 6750, loss = 0.1968  
Step 6800, loss = 0.1973  
Step 6850, loss = 0.1956  
Step 6900, loss = 0.1925  
Step 6950, loss = 0.1916  
Step 7000, loss = 0.1991  
Step 7050, loss = 0.1915  
Step 7100, loss = 0.1916  
Step 7150, loss = 0.2028  
Step 7200, loss = 0.1913  
Step 7250, loss = 0.1926  
Step 7300, loss = 0.1916  
Step 7350, loss = 0.1951  
Step 7400, loss = 0.1913  
Step 7450, loss = 0.1920  
Step 7500, loss = 0.1910  
Step 7550, loss = 0.1915  
Step 7600, loss = 0.1945  
Step 7650, loss = 0.1948  
Step 7700, loss = 0.1982  
Step 7750, loss = 0.1926  
Step 7800, loss = 0.1911  
Step 7850, loss = 0.1925  
Step 7900, loss = 0.1906  
Step 7950, loss = 0.1950

Our reference solution produces, but don't worry if yours is not exactly the same.

Full batch: accuracy: 0.9240, f1\_score: 0.9384  
Mini-batch: accuracy: 0.9415, f1\_score: 0.9533

```
In [17]: y_pred_full = predict(X_test, w_full)
y_pred_minibatch = predict(X_test, w_minibatch)

print('Full batch: accuracy: {:.4f}, f1_score: {:.4f}'
      .format(accuracy_score(y_test, y_pred_full), f1_score(y_test, y_pred_full)))
print('Mini-batch: accuracy: {:.4f}, f1_score: {:.4f}'
      .format(accuracy_score(y_test, y_pred_minibatch), f1_score(y_test, y_pred_min
```

Full batch: accuracy: 0.9240, f1\_score: 0.9384

Mini-batch: accuracy: 0.9415, f1\_score: 0.9533

```
In [18]: plt.figure(figsize=[15, 10])
plt.plot(trace_full, label='Full batch')
plt.plot(trace_minibatch, label='Mini-batch')
plt.xlabel('Iterations * 50')
plt.ylabel('Loss  $\mathcal{L}(\mathbf{w})$ ')
plt.legend()
plt.show()
```

