

Programming assignment 4: SVM

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.datasets import make_blobs

from cvxopt import matrix, solvers
```

Your task

In this sheet we will implement a simple binary SVM classifier. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any `numpy` functions. No other libraries / imports are allowed.

To solve optimization tasks we will use `CVXOPT` <http://cvxopt.org/> - a Python library for convex optimization. If you use `Anaconda`, you can install it using


```
conda install -c conda-forge cvxopt
```

Generate and visualize the data

```
In [3]: N = 200 # number of samples
D = 2 # number of dimensions
C = 2 # number of classes
seed = 1234 # for reproducible experiments

alpha_tol = 1e-4 # threshold for choosing support vectors

X, y = make_blobs(n_samples=N, n_features=D, centers=C, random_state=seed)
y[y == 0] = -1 # it is more convenient to have {-1, 1} as class labels (instead of
y = y.astype(float)
plt.figure(figsize=[10, 8])
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.show()
```



Task 1: Solving the SVM dual problem

Remember, that the SVM dual problem can be formulated as a Quadratic programming (QP) problem. We will solve it using a QP solver from the `CVXOPT` library.

We use the following form of a QP problem:

$$\begin{aligned} & \underset{\mathbf{x}}{\text{minimize}} && \frac{1}{2} \mathbf{x}^T \mathbf{P} \mathbf{x} + \mathbf{q}^T \mathbf{x} \\ & \text{subject to} && \mathbf{G} \mathbf{x} \leq \mathbf{h} \text{ and } \mathbf{A} \mathbf{x} = \mathbf{b}. \end{aligned}$$

Your task is to formulate the SVM dual problems as a QP of this form and solve it using `CVXOPT`, i.e. specify the matrices \mathbf{P} , \mathbf{G} , \mathbf{A} and vectors \mathbf{q} , \mathbf{h} , \mathbf{b} .

```
In [5]: def solve_dual_svm(X, y):
        """Solve the dual formulation of the SVM problem.

        Parameters
        -----
        X : array, shape [N, D]
            Input features.
        y : array, shape [N]
            Binary class labels (in {-1, 1} format).
```

```

Returns
-----
alphas : array, shape [N]
    Solution of the dual problem.
"""
### TODO: Your code below ###
# These variables have to be of type cvxopt.matrix
# Compute the kernel matrix
N, D = X.shape
P = matrix(np.einsum("i,j,ik,jk->ij", y, y, X, X))
q = matrix(-np.ones([N, 1]))
G = matrix(-np.eye(N))
h = matrix(np.zeros(N))
A = matrix(y.reshape(1, -1))
b = matrix(np.zeros(1))
solvers.options['show_progress'] = False
solution = solvers.qp(P, q, G, h, A, b)
alphas = np.array(solution['x'])
return alphas.reshape(-1)

```

Task 2: Recovering the weights and the bias

```

In [7]: def compute_weights_and_bias(alpha, X, y):
    """Recover the weights w and the bias b using the dual solution alpha.

    Parameters
    -----
    alpha : array, shape [N]
        Solution of the dual problem.
    X : array, shape [N, D]
        Input features.
    y : array, shape [N]
        Binary class labels (in {-1, 1} format).

    Returns
    -----
    w : array, shape [D]
        Weight vector.
    b : float
        Bias term.
    """
    ### TODO: Your code below ###
    # Compute weight vector w
    w = np.dot(X.T, alpha * y)

    # Identify support vectors
    support_vector_indices = alpha > alpha_tol
    support_vectors = X[support_vector_indices]
    support_labels = y[support_vector_indices]
    support_alphas = alpha[support_vector_indices]

    # Compute bias using support vectors
    biases = support_labels - np.dot(support_vectors, w)

```

```
b = np.sum(support_alphas * biases) / np.sum(support_alphas) # Weighted mean f
return w, b
```

Visualize the result (nothing to do here)

```
In [8]: def plot_data_with_hyperplane_and_support_vectors(X, y, alpha, w, b):
        """Plot the data as a scatter plot together with the separating hyperplane.

        Parameters
        -----
        X : array, shape [N, D]
            Input features.
        y : array, shape [N]
            Binary class labels (in {-1, 1} format).
        alpha : array, shape [N]
            Solution of the dual problem.
        w : array, shape [D]
            Weight vector.
        b : float
            Bias term.
        """
        plt.figure(figsize=[10, 8])
        # Plot the hyperplane
        slope = -w[0] / w[1]
        intercept = -b / w[1]
        x = np.linspace(X[:, 0].min(), X[:, 0].max())
        plt.plot(x, x * slope + intercept, 'k-', label='decision boundary')
        plt.plot(x, x * slope + intercept - 1/w[1], 'k--')
        plt.plot(x, x * slope + intercept + 1/w[1], 'k--')
        # Plot all the datapoints
        plt.scatter(X[:, 0], X[:, 1], c=y)
        # Mark the support vectors
        support_vecs = (alpha > alpha_tol)
        plt.scatter(X[support_vecs, 0], X[support_vecs, 1], c=y[support_vecs], s=250, m
        plt.xlabel('$x_1$')
        plt.ylabel('$x_2$')
        plt.legend(loc='upper left')
```

The reference solution is

```
w = array([0.73935606 0.41780426])
b = 0.919937145
```

Indices of the support vectors are

```
[ 78 134 158]
```

```
In [9]: alpha = solve_dual_svm(X, y)
        w, b = compute_weights_and_bias(alpha, X, y)
        print("w =", w)
```

```
print("b =", b)
print("support vectors:", np.arange(len(alpha))[alpha > alpha_tol])
```

```
w = [0.73935606 0.41780426]
b = 0.9199371344144434
support vectors: [ 78 134 158]
```

```
In [10]: plot_data_with_hyperplane_and_support_vectors(X, y, alpha, w, b)
plt.show()
```

